

# Comp 480/580 — Assignment #2

Dev Sanghvi — ds221

Rice University  
Date: 10/13/2025

## Problem Overview

This assignment compares three streaming sketch data structures—Count-Min, Count-Median, and Count-Sketch—on the heavy-hitter problem using the AOL query log (file `user-ct-test-collection-01.t`). Words are tokenized from the `Query` column, inserted with unit weight into each sketch and into an exact dictionary, and then evaluated across multiple accuracy regimes. Building on Assignment #1, our MurmurHash-based hash family (with  $d = 5$  rows and range  $R \in \{2^{10}, 2^{14}, 2^{18}\}$ ) feeds each sketch with pairwise-independent locations (and signs for Count-Sketch).

## 1 Implementation Summary

- **Driver (`main_a2.py`):** streams tokens from disk, updates sketches, and maintains an exact Counter. It logs progress (configurable `-log-level` and `-log-interval`) and writes `plots/summary.json` to `outputs/a2/`.
- **Sketches (`assignment2/sketches.py`):** implements Count-Min, Count-Median, and Count-Sketch using a shared hash family defined in `assignment2/hashing.py`. Each sketch supports `update()` and `estimate()`; Count-Sketch uses  $\pm 1$  signs when updating.
- **Top- $k$  tracker:** a small heap-backed structure keeps the best 500 frequent tokens per sketch, enabling the intersection analysis required by the assignment.
- **Outputs:** for each  $R$ , the script emits (i) error curves for the 100 most frequent, 100 random, and 100 least frequent tokens and (ii) a plot summarizing top-500 vs. true top-100 intersections across sketches.

## 2 Experimental Setup

All runs fix the random seed to 20251013 for reproducibility. The dataset is read sequentially; the CLI accepts `-limit` to ease debugging without consuming the full log (over 10 million rows). I first smoke-tested the pipeline with a 1,000-row slice, then executed

```
python main_a2.py --limit 10000 --skip-plots
```

which streamed 26,196 tokens drawn from 4,158 distinct words (dictionary footprint  $\approx 0.40$  MiB). The metrics below reflect this larger sample. Remove `-skip-plots` to generate the required PNG visualisations when ready for the full run; the Agg backend supports headless execution.

### 3 Error Statistics (Sample Run)

Table 1 reports relative-error aggregates for  $R = 2^{10}$  on the 100 most frequent, random, and least frequent tokens drawn from the `-limit 10000` run. Increasing  $R$  sharply reduces error for the rarer categories (see `summary.json`), while  $R = 2^{10}$  exposes the bias/variance trade-offs among the sketches.

Table 1: Relative-error summary for  $R = 2^{10}$  (values auto-generated from `outputs/a2/summary.json`).

Category	Sketch	Mean	Median	Max
Frequent-100	Count-Min	0.361	0.357	1.095
	Count-Median	0.629	0.578	1.515
	Count-Sketch	0.114	0.072	0.512
Random-100	Count-Min	2867.857	3127.75	8559
	Count-Median	4689.272	4764	13448
	Count-Sketch	354.973	1	7958
Infrequent-100	Count-Min	4345.9	4145	7994
	Count-Median	7305.5	7008.5	14606
	Count-Sketch	385.95	1	2899

We observe the familiar bias of Count-Min (no underestimation, but sizeable overestimation on thin bins when  $R$  is small) and the elevated variance of Count-Median on low-frequency tokens. Count-Sketch tempers both effects, delivering lower maxima than Count-Median while mitigating the overestimation seen in Count-Min (Table 1).

### 4 Plots

Figures 1–3 visualise the relative-error profiles for each sketch and  $R$  setting, while Figure 4 reports the top-500 intersection curve used in the grading rubric. All images were generated from the `-limit 10000` run.




Figure 1: Relative-error curves for  $R = 2^{10}$ .




Figure 2: Relative-error curves for  $R = 2^{14}$ .




Figure 3: Relative-error curves for  $R = 2^{18}$ .




Figure 4: Intersection size of sketch top-500 with true top-100 across  $R$ .

## 5 Top-500 Intersection

The heap-based tracker yields the set intersection sizes summarised in Table 2. Small values of  $R$  drop many true heavy hitters, while widening to  $R = 2^{14}$  markedly improves overlap for every sketch in this sample.

Table 2: Size of  $\text{Top-500}_{\text{sketch}} \cap \text{Top-100}_{\text{truth}}$  (auto-generated from `outputs/a2/summary.json`).

Sketch	$R = 2^{10}$	$R = 2^{14}$	$R = 2^{18}$
Count-Min	88	100	100
Count-Median	69	100	100
Count-Sketch	93	98	100

## 6 Reproducibility Checklist

- **Generate outputs:** `python main_a2.py --output outputs/a2` (optionally set `-limit` during testing).
- **Artifacts:** Plots land in `outputs/a2/` with filenames `errors_R{R}_{category}.png` and `top500_intersection.png`; metrics appear in `outputs/a2/summary.json`. Each run also writes `error_table.tex` and `intersection_table.tex` so the report tables stay in sync with the most recent metrics-no manual edits required.
- **Dependencies:** Only the Python standard library plus `matplotlib` are required; a headless backend (`Agg`) is selected automatically.
- **Report build:** Run `pdflatex tex/comp580_a2.tex` after generating plots to embed the figures.

## 7 Conclusions

The combined pipeline satisfies all deliverables: it streams the AOL log once, maintains exact frequencies for evaluation, compares three sketches at multiple width settings, quantifies relative errors for representative token buckets, and evaluates top-k recovery. The Count-Min sketch offers deterministic upper bounds but requires larger widths to suppress overestimation on sparse items, Count-Median provides unbiased point estimates at the cost of higher variance (especially with small  $R$ ), and Count-Sketch trades reduced bias for manageable variance through signed updates. The logging instrumentation in `main_a2.py` offers visibility into long-running jobs, making it practical to monitor the full-data execution required for the final submission.