

# Introduction to Stream Computing and Reservoir Sampling

Agniva Chowdhury

Lecture 7: COMP 480/580  
09/18/2025

# Data streams

- Data that are continuously generated by many sources at very fast rates.
- Examples:
  - Google queries
  - Twitter feeds
  - Financial markets
  - Internet traffic
- We do not have complete information (e.g., size) on the entire dataset.
- Convenient to think about data as infinite.

**How do you make critical calculations about the stream using limited amount of memory?**

# Applications

- Mining query streams:
  - Google wants to know what queries are more frequent today than yesterday.
- Mining click streams
  - Yahoo wants to know which of its pages are getting an unusual number of hits in the past hour.
- Mining social network news feeds:
  - Such as trending topics on Twitter, Facebook, etc.

# Applications (cont'd)

- Sensor networks:
  - Many sensors feeding into a central controller.
- Telephone call records:
  - Data feeds into customer bills as well as settlements between telephone companies.
- IP packets monitored at a switch:
  - Gather information for optimal routing.
  - Detect denial-of-service attacks.

# Motivations

- In massive or infinite data streams, we can't store all data.
- Still, we may need to:
  - Estimate statistics (e.g., mean, median).
  - Build models (e.g., classifiers).
  - Detect trends or anomalies.
- A uniform random sample:
  - Provides an unbiased summary of the full stream.
  - Can be used to approximate global properties.

# Sampling

**Random sampling is a powerful and general tool in data analysis. We'll see several variants and applications.**

- Pick a small random set  $S$  from a large set.
- Estimate quantity of interest on  $S$  instead of entire data set.
- Analysis relies on sampling strategy, sample size, and estimation algorithm.

Basic sampling strategy: uniform sample of size  $k$  from set of size  $m$

- *with replacement*: pick a uniformly random number  $i \in \{1, 2, \dots, m\}$ .
- *without replacement*: pick a single set uniformly from all sets of
- size  $k$  (with cardinality  $\binom{m}{k}$ ).

# Streaming model

- We are given a stream of objects/items/tokens  $e_1, e_2, e_3 \dots$  arriving one by one.
- The total number of items is unknown in advance.
- Let  $m$  be the number of items seen so far.
- The algorithm has limited memory, enough to store only  $k \ll m$  items.
- We don't know how many points we will observe in advance.
- Want to compute interesting functions over the input data.

# Basic question: how to sample from a stream

How do we pick a single uniform sample without knowing length of stream in advance?

- **Goal:** we want every item in a stream of unknown length to have equal probability of being selected, without storing the entire stream.
- We don't know how to adjust the probability without knowing the total length in advance.

# Our strategy: Reservoir sampling

How do we pick a single ( $k = 1$ ) uniform sample without knowing the length of stream in advance?

## Uniform Sample:

$s \leftarrow \emptyset$

$m \leftarrow 0$

While (stream is not done) :

- $m \leftarrow m + 1$
- $e_m$  is the current item
- Toss a biased coin with  $P(H) = \frac{1}{m}$
- If  $H$ :

$S \leftarrow e_m$

EndWhile

Output  $s$  as the sample

# Reservoir sampling: A key result

Let  $m$  be the number of items in a stream:  $e_1, e_2, e_3 \dots, e_m$ . Let  $s$  be the final output of the reservoir sampling algorithm with  $k = 1$ . Then

$$\mathbb{P}(s = e_j) = \frac{1}{m} \text{ for all } j = 1, 2, \dots, m.$$

Proof:

1. Base case ( $m = 1$ ): Only one item  $e_1$ . The algorithm selects it with probability 1.
2. Suppose it works for  $m - 1$  i.e.,  $\mathbb{P}(s = e_j) = \frac{1}{m-1}$  for each  $j = 1, 2, \dots, m - 1$ .
3. At step  $m$ :
  - The new item  $e_m$  either replaces the current sample with prob.  $1/m$  or the previous sample was kept with prob.  $1 - \frac{1}{m}$
  - For  $j < m$ ,  $\mathbb{P}(s = e_j) = \mathbb{P}(s = e_j \text{ before}) \cdot \mathbb{P}(\text{not replaced at } m^{\text{th}} \text{ step}) = \frac{1}{m-1} \cdot \left(1 - \frac{1}{m}\right) = \frac{1}{m}$
  - For  $j = m$ ,  $\mathbb{P}(s = e_j) = \mathbb{P}(\text{replaced at } m^{\text{th}} \text{ step}) = \frac{1}{m}$

# Reservoir sampling for size $k > 1$

- Want to pick  $k$  samples for  $k > 1$ . How?
- With replacement: Easy. simply run single sample algorithm independently in parallel and store the  $k$  items.
- Without replacement?

# Reservoir sampling for size $k > 1$

## Sampling without replacement:

$S[1 \dots k] \leftarrow \emptyset$

$m \leftarrow 0$

While (stream is not done) :

- $m \leftarrow m + 1$
- $e_m$  is the current item
- If  $m < k$ :  $S[m] \leftarrow e_m$
- Else:

$r \leftarrow$  uniform random number in  $\{1, 2, \dots, m\}$

if  $r \leq k$ :  $S[r] = e_m$

EndWhile

Output  $S$

# Reservoir sampling for size $k > 1$

Let  $m$  be the number of items in a stream:  $e_1, e_2, e_3 \dots, e_m$ . Let  $S$  be the final array of samples maintained by the algorithm with reservoir size  $1 \leq k \leq m$ . Then for any item  $e_j$  with  $1 \leq j \leq m$ ,

$$\mathbb{P}(e_j \in S) = \frac{k}{m}$$

Proof:

1. Base case ( $m = k$ ): each is included with probability 1. Correct since  $\frac{k}{k} = 1$ .
2. Suppose after processing  $m$  items, each has probability  $\frac{k}{m}$  of being in the reservoir.
3. Now, for  $(m + 1)^{th}$  item,  $\mathbb{P}(r \leq k) = \frac{k}{m+1}$
4. For  $j \leq m$ : ??

# Weighted sampling for size $k = 1$

- Uniform sampling treats all items equally.
- But in many applications, some items are more important than others:
  - Items with high frequency or importance.
  - Events with high value, risk, or uncertainty.
  - Weighted sampling gives priority to such important items.
- In streaming settings:
  - We want to sample important items proportionally to their weight.
  - But we can't store the whole stream to normalize weights.

Can we select items with probability proportional to weight  
using **one pass and small memory?**

# Weighted sampling for size $k = 1$

- Each item  $e_j$  in the stream has an associated weight  $w_j > 0$ .
- **Goal:** sample one item with probability proportional to its weight.

## Weighted Sample:

For  $j = 1$  to  $m$ :

- Generate  $U_j \sim \text{Uniform}(0, 1)$
- Compute priority  $r_j = U_j^{1/w_j}$

EndFor

Output: Keep the item with largest  $r_j$

- Due to Efraimidis and Spirakis (2006)

# Weighted sampling for size $k = 1$

Let  $m$  be the number of items in a stream:  $e_1, e_2, e_3 \dots, e_m$ . Each item has weight  $w_j > 0$ . Let  $s$  be the final sample maintained by the algorithm with size = 1. Then for any item  $e_j$  with  $1 \leq j \leq m$ ,

$$\mathbb{P}(s = e_j) = \frac{w_j}{W} \text{ where } W = \sum_{j=1}^m w_j$$

Proof:

# Extensions

- Frequent items / Heavy hitters in streams.
- Sketching & randomized summaries (Count-Min, CountSketch)
- Sublinear-time algorithms for streaming models
- Reservoir Sampling with deletion / time-decay models

⋮  
⋮

**Thank you  
Questions?**