

# COMP 480 Lecture Notes: Consistent Hashing

Scribed by Quan Tran (qmt1), Murad Imre (mi15), Camden Graham (cg5)

September 23, 2025

## 1 Bloom Filters and Web Caching

### 1.1 Motivation for Caching

When a webpage is requested multiple times, it is efficient to cache it.

- Example: If you request `amazon.com` and did so recently, the page can be loaded from your local cache.
- Otherwise, you incur a **cache miss**, requiring a download from the remote server.

### 1.2 Benefits

- Faster user experience: cached requests load immediately.
- System-level improvements: reduced network traffic, less congestion, fewer communication overheads, fewer dropped packets.

### 1.3 Rule of Thumb

**If something is used more than once, cache it.** This heuristic mirrors how humans optimize repeated tasks.

### 1.4 Examples

- Word usage: A few high-frequency words (“head” of the distribution) account for most traffic. Handling 80–90% of traffic is possible with caching.
- Netflix: Trending movies are cached near users, loading quickly. Rare movies are slower because they may not be cached.
- Geography: Requests routed to faraway caches (e.g., Australia) perform poorly compared to nearby caches.

### 1.5 Shared Caches

- Idea: A shared cache for multiple nearby users (e.g., all Rice students).
- Benefits: More aggregate hits, less latency.




Figure 1: Word frequency distribution illustrating cache efficiency.

- Example: Akamai Technologies (founded 1998, valued at \$10B+) built a CDN business around this idea.

## 1.6 Challenges

- Storing recently accessed pages for many users requires massive, fast memory.
- Akamai's plan: perform this in main memory for speed.
- At large scale, caches must be distributed over many machines.
- Question: If there are 100 cache servers, which one should handle `amazon.com`?

## 2 Naïve Solutions and Their Limits

### 2.1 Initial Ideas

- Poll all servers to check if they hold a cached copy — infeasible.
- Use hashing: map `amazon.com` to server  $h(\text{amazon.com})$ .

### 2.2 Problem: Dynamic Changes

- If machines are added/removed frequently, modulo hashing ( $h(x) = x \bmod n$ ) forces nearly all keys to remap.
- Example: Switching from  $\bmod 12$  to  $\bmod 13$  shifts almost everything.

### 2.3 Machine Failures at Scale

Jeff Dean's observations of Google data centers:

“In each cluster’s first year,  $\sim 1000$  machines fail; thousands of disks fail; power units crash; 20 racks go offline; overheating can shut down most servers within 5 minutes.”

This highlights the importance of robustness.

## 3 Consistent Hashing

### 3.1 Definition

- Hash both servers and objects into the same numeric space, visualized as a **ring**.
- Rule: Assign object  $x$  to the nearest server in the clockwise direction from  $h(x)$ .




Figure 2: Consistent hashing ring showing server placement and object assignment.

### 3.2 Advantages

- Adding a server only requires moving the data that maps to it.
- Removing a server reroutes its objects to the next clockwise server.
- Only a small fraction of objects move, unlike in standard modulo hashing.

### 3.3 Search Time

- Naïve: scanning clockwise —  $O(n)$ .
- Improved: Store server hashes in a balanced binary search tree; lookup becomes  $O(\log n)$ .

### 3.4 Historical Timeline

- 1997: First paper on consistent hashing presented at STOC. Originally rejected for being “impractical.”
- 1998: Akamai founded.
- 1999: Star Wars trailer crashes Apple’s site; Akamai’s CDN handles the demand.

- April 1, 1999: Steve Jobs calls Akamai's President Paul Sagan. Sagan thinks it's an April Fool's prank and hangs up.
- September 11, 2001: Co-founder Danny Lewin tragically killed on the first plane that hit the World Trade Center.
- 2001: Repurposed for peer-to-peer (P2P) networks — foundation of distributed hash tables (DHTs).
- 2006: Adopted by Amazon Dynamo for scalable, memory-based storage on commodity hardware.

## 4 Mathematical Analysis of Load Balancing

### 4.1 Expected Load

- Given  $m$  items and  $n$  machines, the expected load of each machine is  $\frac{m}{n}$ .
  - This follows by symmetry: each machine is equally likely to be assigned an item.
- When a machine is added, the expected number of items that move to the new machine is  $\frac{m}{n+1}$ .
  - Again, this is a straightforward argument by symmetry.
- With high probability, no machine owns more than  $O\left(\frac{\log n}{n}\right)$  fraction of the total items.

### 4.2 Proof (from lecture)

- With high probability, no machine is severely **overloaded**.
- Question: Can we also guarantee that no machine is **underloaded**?
- Analogy: The **birthday paradox** illustrates how random allocation can still produce surprising gaps or collisions, raising the underloaded concern.

**Overload bound (interval method):**

- Fix an interval  $I$  of length  $\frac{2\log n}{n}$ .
- Probability that no machine lands in this interval is:

$$\left(1 - \frac{2\log n}{n}\right)^n \approx e^{-2\log n} = \frac{1}{n^2}.$$

- Divide the unit circle into  $\frac{n}{2\log n}$  disjoint intervals.
  - By the union bound, the probability that *any* interval is empty is at most:
- $$\frac{n}{2\log n} \cdot \frac{1}{n^2} \leq \frac{1}{n}.$$
- Therefore, with probability at least  $1 - \frac{1}{n}$ , every interval has at least one machine.

- Thus each machine controls at most  $\frac{4\log n}{n}$  fraction of the load, with high probability.

#### Variance reduction (replication):

- To balance underloaded cases, create  $K$  **virtual copies** of each machine.
- The load on each machine is then modeled as the sum of  $K$  i.i.d. random variables.
- By standard concentration bounds, the sum is sharply concentrated around its mean.
- This ensures workloads are balanced across machines with high probability.